

Alzheimer's (DementiaBank Pitt): disease information from the audio encoder through the language model
Qwen2-Audio, per-stage linear probe, 5-fold GroupKFold by speaker, 468 cookie-task segments, 228 speakers

