

qwen2 audio encoder, one linear probe per layer, all 33 layers
within dataset, whole speakers held out (5 fold GroupKFold), audio only, no transcript, 30 s windows

