

qwen2 audio encoder, one linear probe per layer, all 33 layers
 within dataset, whole speakers held out, audio only, no transcript, 30 s windows

