

qwen2 audio encoder, a probe trained on every layer
within dataset, whole speakers held out, audio only, no transcript, 30 s windows

