



Probe AUC through the whole model, qwen2 audio
encoder layers 0 to 32, then the projector (shaded), then LLM layers at the audio token positions


KCL read, Parkinson's (37 speakers)




Neurovoz read, Parkinson's (107 speakers)



DAIC interview, depression (275 speakers)



Pitt cookie task, Alzheimer's (228 speakers)



stage