


Probe AUC through the whole model, qwen2 audio: encoder layers, then the projector (shaded), then LLM layers at the audio token positions

KCL read (Parkinson's, English, 37 clips, 37 speakers)



DementiaBank Pitt cookie task (Alzheimer's, English, 468 clips, 228 speakers)

