

alzheimers, qwen2-audio, same 468 clips: with audio the answer lands where the transcript alone lands, so it was reading the words. the trained versions use the audio

